

```
In [1]: # Predict asteroid diameter values using 'Asteroid.csv' dataset from Kaggle co
ntributed by Victor Basu
        # Link: https://www.kaggle.com/basu369victor/prediction-of-asteroid-diamet
er
        # Model: XGBRegressor
        # Notes on data:
            # data is medium size comprising of 839736 entries and 27 columns
            # for a small portion of the data (~ 1/6) the asteroids diameters are know
n -
            # this portion will be used to train and validate the model
            # subsequently the model will be used to predict the diameters for the dat
a in which this information is missing

        # Essential updates (6/2020) from previous project version (2/2020)
            # Improvements in data processing and data visualization
            # Comparison between XGBRegressor model and Linear Regression model is dis
carded -->
            # XGBRegressor model optimization via hyperparameter tuning is added i
nstead
            # Statistics of residuals - distribution, mean and standard deviation - re
place absolute error statistics
            # as model performance metrics
```

```
In [2]: # Import Libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set(style = "whitegrid", font_scale = 1.5)
```

```
In [3]: # Ignore warnings

import warnings
warnings.filterwarnings('ignore')
```

In [4]: *# Read data*

```
data = pd.read_csv('Asteroid.csv', low_memory = False)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839736 entries, 0 to 839735
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   full_name              839736 non-null object
1   a                      839734 non-null float64
2   e                      839736 non-null float64
3   G                      119 non-null   float64
4   i                      839736 non-null float64
5   om                    839736 non-null float64
6   w                      839736 non-null float64
7   q                      839736 non-null float64
8   ad                    839730 non-null float64
9   per_y                 839735 non-null float64
10  data_arc              823947 non-null float64
11  condition_code        838743 non-null object
12  n_obs_used            839736 non-null int64
13  H                     837042 non-null float64
14  diameter              137681 non-null object
15  extent                18 non-null   object
16  albedo                136452 non-null float64
17  rot_per               18796 non-null float64
18  GM                    14 non-null   float64
19  BV                    1021 non-null float64
20  UB                     979 non-null  float64
21  IR                     1 non-null    float64
22  spec_B                1666 non-null object
23  spec_T                980 non-null  object
24  neo                   839730 non-null object
25  pha                   822814 non-null object
26  moid                  822814 non-null float64
dtypes: float64(18), int64(1), object(8)
memory usage: 173.0+ MB
```

In [5]: *# Print data column names for use in the code below*

```
data.columns
```

```
Out[5]: Index(['full_name', 'a', 'e', 'G', 'i', 'om', 'w', 'q', 'ad', 'per_y',
              'data_arc', 'condition_code', 'n_obs_used', 'H', 'diameter', 'extent',
              'albedo', 'rot_per', 'GM', 'BV', 'UB', 'IR', 'spec_B', 'spec_T', 'neo',
              'pha', 'moid'],
              dtype='object')
```

```
In [6]: # Select only features with meaningful amount of non-null values -->
        # drop 'G', 'extent', 'GM', 'BV', 'UB', 'IR', 'spec_B', and 'spec_T'
        # In addition, drop 'full_name' and 'n_obs_used' which are not meaningful for
        # the problem
        # Place the target 'diameter' at the end for easier separation of features, X,
        # and target, y, later on

data = data[['a', 'e', 'i', 'om', 'w', 'q', 'ad', 'per_y', 'data_arc', 'condition_code',
            'H', 'albedo', 'neo', 'pha', 'moid', 'diameter']]
data.head(10)
```

Out[6]:

	a	e	i	om	w	q	ad	per_y	data_ar
0	2.769165	0.076009	10.594067	80.305532	73.597694	2.558684	2.979647	4.608202	8822.
1	2.772466	0.230337	34.836234	173.080063	310.048857	2.133865	3.411067	4.616444	72318.
2	2.669150	0.256942	12.988919	169.852760	248.138626	1.983332	3.354967	4.360814	72684.
3	2.361418	0.088721	7.141771	103.810804	150.728541	2.151909	2.570926	3.628837	24288.
4	2.574249	0.191095	5.366988	141.576604	358.687608	2.082324	3.066174	4.130323	63431.
5	2.425160	0.203007	14.737901	138.640203	239.807490	1.932835	2.917485	3.776755	62329.
6	2.385334	0.231206	5.523651	259.563231	145.265106	1.833831	2.936837	3.684105	62452.
7	2.201764	0.156499	5.886955	110.889330	285.287462	1.857190	2.546339	3.267115	62655.
8	2.385637	0.123114	5.576816	68.908577	6.417369	2.091931	2.679342	3.684806	61821.
9	3.141539	0.112461	3.831560	283.202167	312.315206	2.788240	3.494839	5.568291	62175.

```
In [7]: # 1) Data Processing and EDA
```


In [8]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839736 entries, 0 to 839735
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   a                   839734 non-null  float64
1   e                   839736 non-null  float64
2   i                   839736 non-null  float64
3   om                  839736 non-null  float64
4   w                   839736 non-null  float64
5   q                   839736 non-null  float64
6   ad                  839730 non-null  float64
7   per_y              839735 non-null  float64
8   data_arc           823947 non-null  float64
9   condition_code     838743 non-null  object
10  H                   837042 non-null  float64
11  albedo             136452 non-null  float64
12  neo                839730 non-null  object
13  pha                822814 non-null  object
14  moid               822814 non-null  float64
15  diameter           137681 non-null  object
dtypes: float64(12), object(4)
memory usage: 102.5+ MB
```

In [9]: *# Features 'condition_code', 'neo', and 'pha' appear to be categorical --> examine these features*

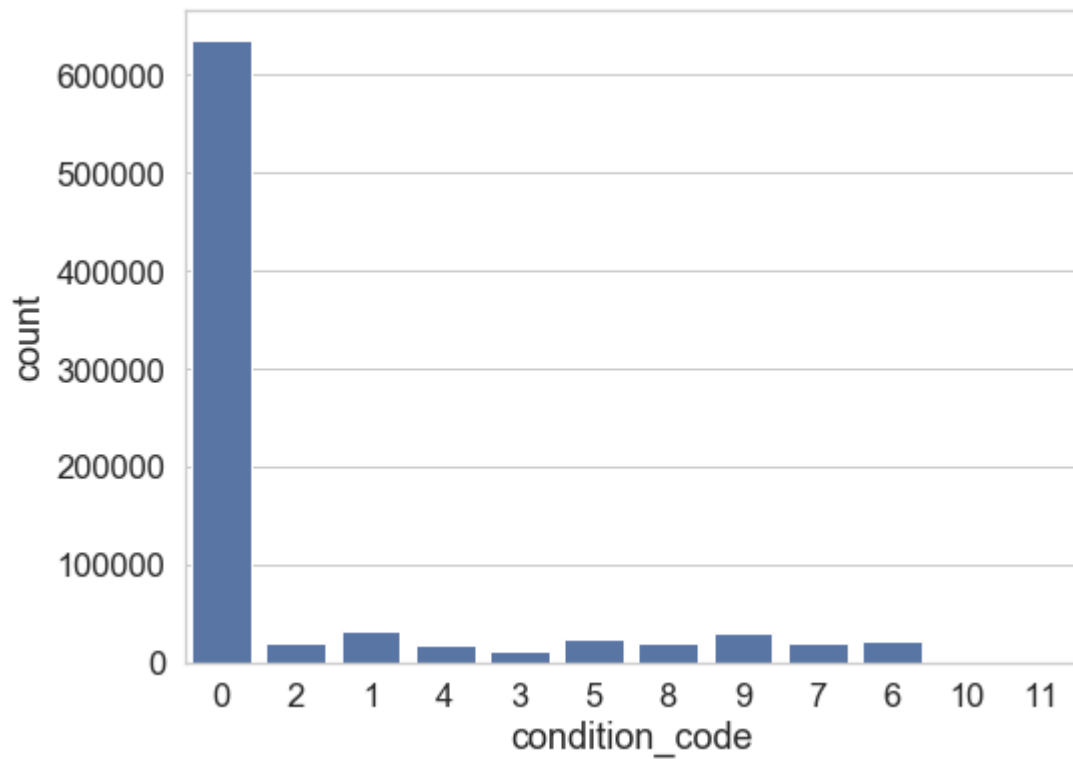
```
In [10]: # Examine 'condition_code'

plt.figure(figsize = (8, 6))
sns.countplot(data['condition_code'], color = 'b')
plt.show()
```




```
In [11]: # This is a categorical feature with majority of data points having values = 0
# Note that 'condition_code' values includes both numbers and letters
```

```
In [12]: # Assign numeric values to the categorical values 'E' and 'D'  
  
data['condition_code'].replace({'E': 10, 'D': 11}, inplace=True)  
  
plt.figure(figsize = (8, 6))  
sns.countplot(data['condition_code'], color = 'b')  
plt.show()
```



```
In [13]: # Examine 'neo'

plt.figure(figsize = (8, 6))
sns.countplot(data['neo'])
plt.show()
```




```
In [14]: # Categorical feature --> majority of data points = N
```

```
In [15]: # Replace categorical values, N and Y, with numerical values of 0 and 1, respectively

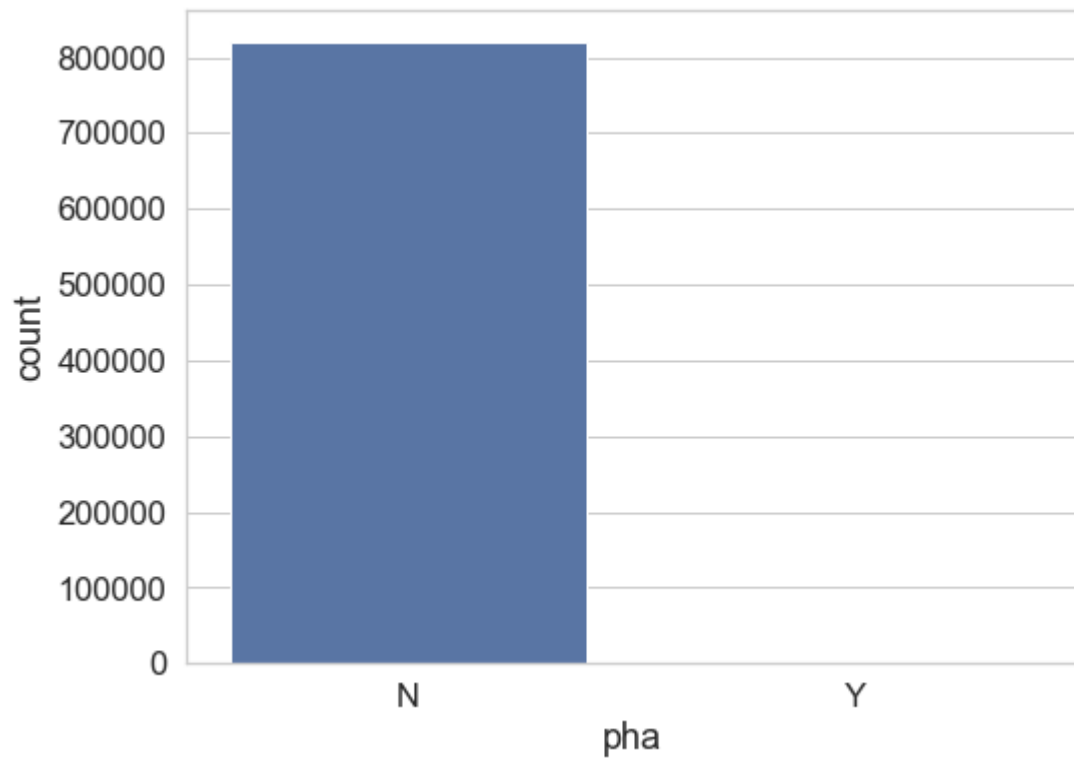
data['neo'].replace({'N': 0, 'Y': 1}, inplace=True)

plt.figure(figsize = (8, 6))
sns.countplot(data['neo'], color = 'b')
plt.show()
```

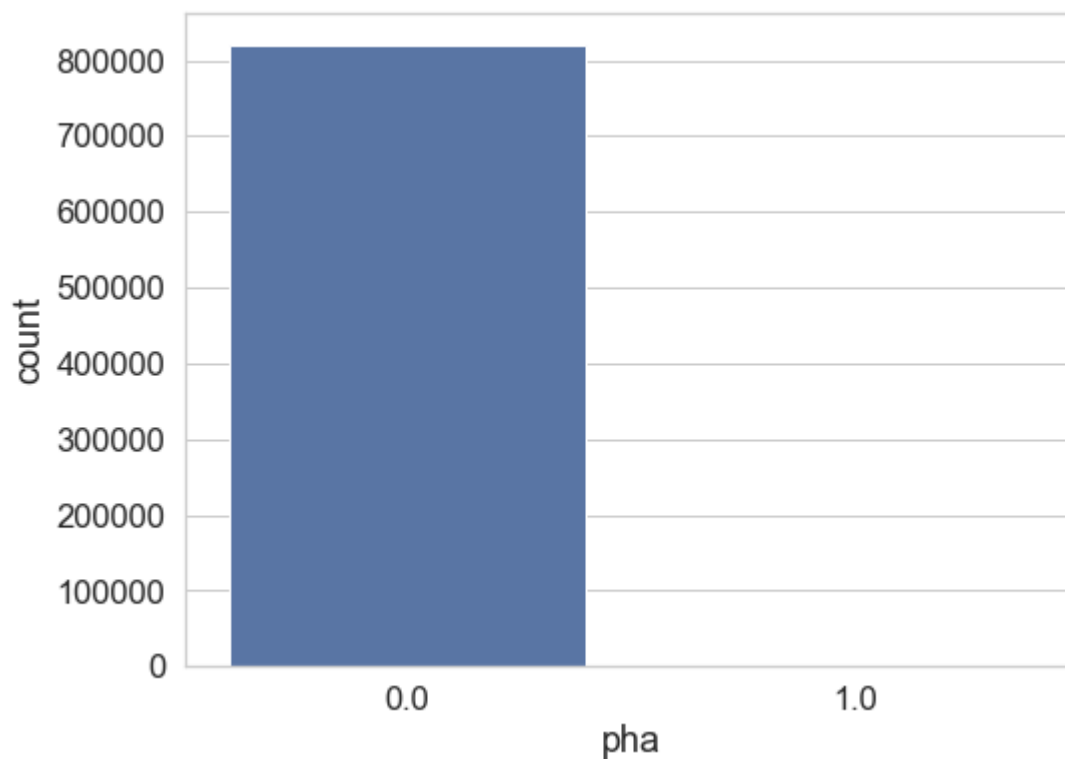



```
In [16]: # Examine 'pha'

plt.figure(figsize = (8, 6))
sns.countplot(data['pha'])
plt.show()
```



```
In [17]: # Categorical feature --> majority of data points = N  
  
# Replace the categorical values, N and Y, with numerical values of 0 and 1, r  
# respectively  
  
data['pha'].replace({'N': 0, 'Y': 1}, inplace=True)  
  
plt.figure(figsize = (8, 6))  
sns.countplot(data['pha'], color = 'b')  
plt.show()
```



In [18]: `# Examine target, 'diameter'`

`data.head(10)`

Out[18]:

	a	e	i	om	w	q	ad	per_y	data_ar
0	2.769165	0.076009	10.594067	80.305532	73.597694	2.558684	2.979647	4.608202	8822.
1	2.772466	0.230337	34.836234	173.080063	310.048857	2.133865	3.411067	4.616444	72318.
2	2.669150	0.256942	12.988919	169.852760	248.138626	1.983332	3.354967	4.360814	72684.
3	2.361418	0.088721	7.141771	103.810804	150.728541	2.151909	2.570926	3.628837	24288.
4	2.574249	0.191095	5.366988	141.576604	358.687608	2.082324	3.066174	4.130323	63431.
5	2.425160	0.203007	14.737901	138.640203	239.807490	1.932835	2.917485	3.776755	62329.
6	2.385334	0.231206	5.523651	259.563231	145.265106	1.833831	2.936837	3.684105	62452.
7	2.201764	0.156499	5.886955	110.889330	285.287462	1.857190	2.546339	3.267115	62655.
8	2.385637	0.123114	5.576816	68.908577	6.417369	2.091931	2.679342	3.684806	61821.
9	3.141539	0.112461	3.831560	283.202167	312.315206	2.788240	3.494839	5.568291	62175.

In [19]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839736 entries, 0 to 839735
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   a                839734 non-null  float64
1   e                839736 non-null  float64
2   i                839736 non-null  float64
3   om               839736 non-null  float64
4   w                839736 non-null  float64
5   q                839736 non-null  float64
6   ad               839730 non-null  float64
7   per_y            839735 non-null  float64
8   data_arc         823947 non-null  float64
9   condition_code   838743 non-null  object
10  H                837042 non-null  float64
11  albedo           136452 non-null  float64
12  neo              839730 non-null  float64
13  pha              822814 non-null  float64
14  moid             822814 non-null  float64
15  diameter         137681 non-null  object
dtypes: float64(14), object(2)
memory usage: 102.5+ MB
```

```
In [20]: # Columns 'diameter' and 'albedo' have only about 1/6 of non-null values compared to other features
# Although 'diameter' has numerical values in the table, it appears that it is in string format - data type 'object'
# Convert data to numeric format 'float64'

data = data.astype('float64')

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839736 entries, 0 to 839735
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      839734 non-null float64
1   e                      839736 non-null float64
2   i                      839736 non-null float64
3   om                    839736 non-null float64
4   w                     839736 non-null float64
5   q                     839736 non-null float64
6   ad                    839730 non-null float64
7   per_y                 839735 non-null float64
8   data_arc              823947 non-null float64
9   condition_code        838743 non-null float64
10  H                      837042 non-null float64
11  albedo                136452 non-null float64
12  neo                   839730 non-null float64
13  pha                   822814 non-null float64
14  moid                  822814 non-null float64
15  diameter              137681 non-null float64
dtypes: float64(16)
memory usage: 102.5 MB
```

In [21]: *# Replace all missing values with 0 which is the sparse value expected by XGBoost*

```
data.fillna(0, inplace = True)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 839736 entries, 0 to 839735
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	a	839736 non-null	float64
1	e	839736 non-null	float64
2	i	839736 non-null	float64
3	om	839736 non-null	float64
4	w	839736 non-null	float64
5	q	839736 non-null	float64
6	ad	839736 non-null	float64
7	per_y	839736 non-null	float64
8	data_arc	839736 non-null	float64
9	condition_code	839736 non-null	float64
10	H	839736 non-null	float64
11	albedo	839736 non-null	float64
12	neo	839736 non-null	float64
13	pha	839736 non-null	float64
14	moid	839736 non-null	float64
15	diameter	839736 non-null	float64

```
dtypes: float64(16)
```

```
memory usage: 102.5 MB
```

```
In [22]: # Create data set, data_1, where diameter is known

data_1 = data[data['diameter'] > 0] # values greater than 0 correspond to data
with known diameter
data_1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 137681 entries, 0 to 810411
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      137681 non-null float64
1   e                      137681 non-null float64
2   i                      137681 non-null float64
3   om                    137681 non-null float64
4   w                      137681 non-null float64
5   q                      137681 non-null float64
6   ad                    137681 non-null float64
7   per_y                 137681 non-null float64
8   data_arc              137681 non-null float64
9   condition_code        137681 non-null float64
10  H                      137681 non-null float64
11  albedo                 137681 non-null float64
12  neo                    137681 non-null float64
13  pha                    137681 non-null float64
14  moid                   137681 non-null float64
15  diameter               137681 non-null float64
dtypes: float64(16)
memory usage: 17.9 MB
```

```
In [23]: # Data with known asteroid diameter have total of 137681 entries
```

```
In [24]: # Check data_1

data_1.head(10)
```

Out[24]:

	a	e	i	om	w	q	ad	per_y	data_ar
0	2.769165	0.076009	10.594067	80.305532	73.597694	2.558684	2.979647	4.608202	8822.
1	2.772466	0.230337	34.836234	173.080063	310.048857	2.133865	3.411067	4.616444	72318.
2	2.669150	0.256942	12.988919	169.852760	248.138626	1.983332	3.354967	4.360814	72684.
3	2.361418	0.088721	7.141771	103.810804	150.728541	2.151909	2.570926	3.628837	24288.
4	2.574249	0.191095	5.366988	141.576604	358.687608	2.082324	3.066174	4.130323	63431.
5	2.425160	0.203007	14.737901	138.640203	239.807490	1.932835	2.917485	3.776755	62329.
6	2.385334	0.231206	5.523651	259.563231	145.265106	1.833831	2.936837	3.684105	62452.
7	2.201764	0.156499	5.886955	110.889330	285.287462	1.857190	2.546339	3.267115	62655.
8	2.385637	0.123114	5.576816	68.908577	6.417369	2.091931	2.679342	3.684806	61821.
9	3.141539	0.112461	3.831560	283.202167	312.315206	2.788240	3.494839	5.568291	62175.

In [25]: *# Everything looks fine*

In [26]: *# Create dataset, data_2, where diameter is unknown*

```
data_2 = data[data['diameter'] < data_1['diameter'].min()] # this leaves only
0s which correspond to unknown diameter
data_2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 702055 entries, 681 to 839735
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      702055 non-null float64
1   e                      702055 non-null float64
2   i                      702055 non-null float64
3   om                    702055 non-null float64
4   w                    702055 non-null float64
5   q                    702055 non-null float64
6   ad                   702055 non-null float64
7   per_y               702055 non-null float64
8   data_arc            702055 non-null float64
9   condition_code      702055 non-null float64
10  H                   702055 non-null float64
11  albedo              702055 non-null float64
12  neo                 702055 non-null float64
13  pha                 702055 non-null float64
14  moid                702055 non-null float64
15  diameter            702055 non-null float64
dtypes: float64(16)
memory usage: 91.1 MB
```

In [27]: *# Data with unknown asteroid diameter have total of 702055 entries (more than 5 x that of data_1)*

In [28]: *# Check data_2*

```
data_2.head(10)
```


Out[28]:

	a	e	i	om	w	q	ad	per_y	data
681	2.654040	0.171983	11.505648	190.799958	104.993826	2.197591	3.110489	4.323837	400
698	2.610998	0.410284	15.299180	242.551766	91.399514	1.539746	3.682249	4.219081	425
718	2.638780	0.546301	11.564845	183.887287	156.163668	1.197212	4.080348	4.286601	394
729	2.243362	0.177505	4.234895	95.073806	123.549777	1.845154	2.641570	3.360139	397
842	2.279598	0.209766	7.997717	4.071363	316.957206	1.801415	2.757780	3.441878	375
961	2.908998	0.097329	2.602636	145.481660	223.473847	2.625868	3.192128	4.961619	374
984	2.299979	0.277462	4.056565	290.307048	59.553605	1.661822	2.938137	3.488142	353
1008	2.625175	0.455500	15.769676	229.461495	186.428747	1.429408	3.820942	4.253492	349
1010	2.391976	0.350864	5.494744	132.525452	353.279770	1.552718	3.231235	3.699504	349
1064	2.360276	0.297141	8.362855	330.324142	353.652287	1.658942	3.061610	3.626205	338

In [29]: *# It appears 'albedo' is also unknown in data_2 --> only 0s are shown in table*

Check by plotting data_2['albedo']

```
plt.figure(figsize = (15, 8))
plt.scatter(np.arange(1, len(data_2) + 1), data_2['albedo'], s = 50, c = 'b')
plt.xlabel('Data Point #', fontsize = 20, labelpad = 15)
plt.ylabel('Data_2: albedo', fontsize = 20, labelpad = 15)
plt.show()
```




```
In [30]: # Indeed, almost all 'albedo' data points in data_2 are 0s and it cannot be used in predictions -->
         # remove 'albedo' from both data_1 and data_2
```

```
In [31]: data_1.columns
```

```
Out[31]: Index(['a', 'e', 'i', 'om', 'w', 'q', 'ad', 'per_y', 'data_arc',
               'condition_code', 'H', 'albedo', 'neo', 'pha', 'moid', 'diameter'],
              dtype='object')
```

```
In [32]: # Keep all features in data_1 except 'albedo'

data_1 = data_1[['a', 'e', 'i', 'om', 'w', 'q', 'ad', 'per_y', 'data_arc',
                 'condition_code', 'H', 'neo', 'pha', 'moid', 'diameter']]
data_1.head(10)
```

```
Out[32]:
```

	a	e	i	om	w	q	ad	per_y	data_ar
0	2.769165	0.076009	10.594067	80.305532	73.597694	2.558684	2.979647	4.608202	8822.
1	2.772466	0.230337	34.836234	173.080063	310.048857	2.133865	3.411067	4.616444	72318.
2	2.669150	0.256942	12.988919	169.852760	248.138626	1.983332	3.354967	4.360814	72684.
3	2.361418	0.088721	7.141771	103.810804	150.728541	2.151909	2.570926	3.628837	24288.
4	2.574249	0.191095	5.366988	141.576604	358.687608	2.082324	3.066174	4.130323	63431.
5	2.425160	0.203007	14.737901	138.640203	239.807490	1.932835	2.917485	3.776755	62329.
6	2.385334	0.231206	5.523651	259.563231	145.265106	1.833831	2.936837	3.684105	62452.
7	2.201764	0.156499	5.886955	110.889330	285.287462	1.857190	2.546339	3.267115	62655.
8	2.385637	0.123114	5.576816	68.908577	6.417369	2.091931	2.679342	3.684806	61821.
9	3.141539	0.112461	3.831560	283.202167	312.315206	2.788240	3.494839	5.568291	62175.

```
In [33]: # Keep all features in data_2 except 'albedo' and 'diameter' which is unknown

data_2 = data_2[['a', 'e', 'i', 'om', 'w', 'q', 'ad', 'per_y', 'data_arc', 'condition_code', 'H', 'neo', 'pha', 'moid']]

data_2.head(10)
```

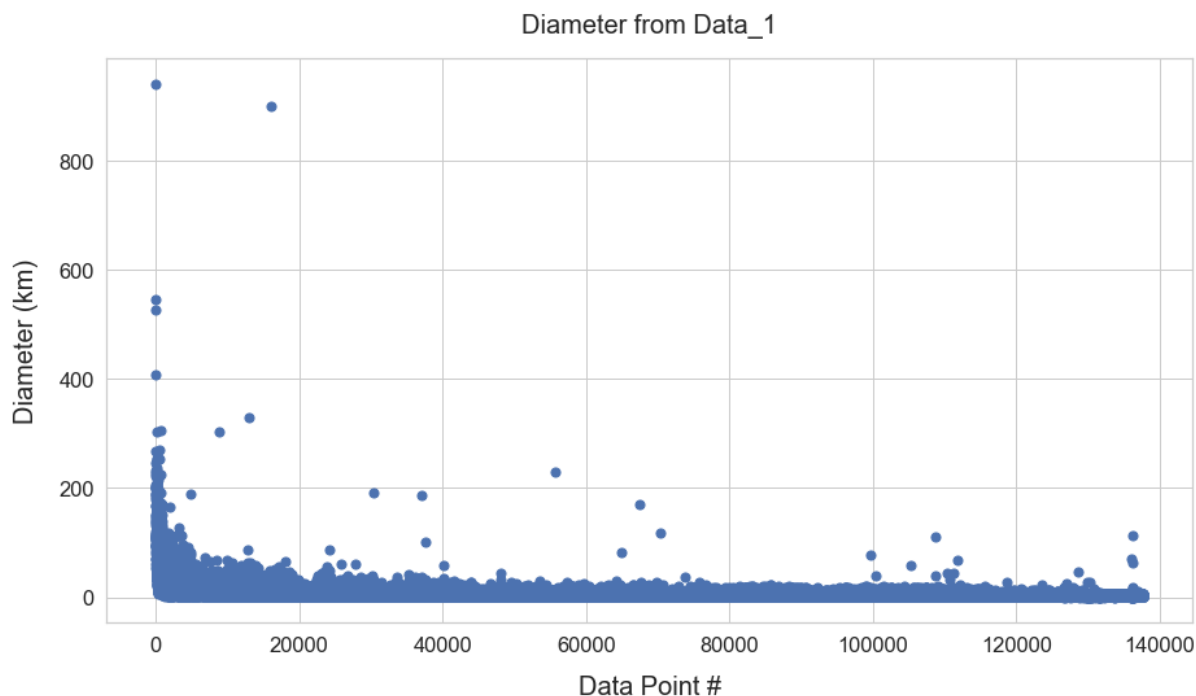
Out[33]:

	a	e	i	om	w	q	ad	per_y	data
681	2.654040	0.171983	11.505648	190.799958	104.993826	2.197591	3.110489	4.323837	400
698	2.610998	0.410284	15.299180	242.551766	91.399514	1.539746	3.682249	4.219081	425
718	2.638780	0.546301	11.564845	183.887287	156.163668	1.197212	4.080348	4.286601	394
729	2.243362	0.177505	4.234895	95.073806	123.549777	1.845154	2.641570	3.360139	397
842	2.279598	0.209766	7.997717	4.071363	316.957206	1.801415	2.757780	3.441878	375
961	2.908998	0.097329	2.602636	145.481660	223.473847	2.625868	3.192128	4.961619	374
984	2.299979	0.277462	4.056565	290.307048	59.553605	1.661822	2.938137	3.488142	353
1008	2.625175	0.455500	15.769676	229.461495	186.428747	1.429408	3.820942	4.253492	349
1010	2.391976	0.350864	5.494744	132.525452	353.279770	1.552718	3.231235	3.699504	349
1064	2.360276	0.297141	8.362855	330.324142	353.652287	1.658942	3.061610	3.626205	338

```
In [34]: # Data_1 has 14 features and target, 'diameter', Left
# Data_2 consists of the same 14 features only -- no 'diameter'
```

```
In [35]: # Visualize 'diameter' from data_1 using scatterplot

plt.figure(figsize = (15, 8))
plt.scatter(np.arange(1, len(data_1) + 1), data_1['diameter'], s = 50, c = 'b'
)
plt.title('Diameter from Data_1', fontsize = 20, pad = 20)
plt.xlabel('Data Point #', fontsize = 20, labelpad = 15)
plt.ylabel('Diameter (km)', fontsize = 20, labelpad = 15)
plt.show()
```



```
In [36]: # It appears 'diameter' has great number of small values and only small portion
of large values
# Get some insights from min, max, median and mean of diameter in data_1
```

```
In [37]: # Min, max, median and mean of diameter in data_1

print("Min diameter in km -->", round(data_1['diameter'].min(), 4))
print("Max diameter in km -->", round(data_1['diameter'].max(), 4))
print("Median diameter in km -->", round(data_1['diameter'].median(), 4))
print("Mean diameter in km -->", round(data_1['diameter'].mean(), 4))
```


```
Min diameter in km --> 0.0025
Max diameter in km --> 939.4
Median diameter in km --> 3.956
Mean diameter in km --> 5.4825
```

```
In [38]: # Key observations:
# 1) Max value is much larger than mean (almost 3 orders of magnitude)
# 2) Despite that, mean and median are very close --> large values are small
portion of the total number of observations
```

```
In [39]: # Explore further by using boxplots
# Important note:
# use log scale due to large spread and disparity between the number of small (majority) and large diameter values
```

```
In [40]: # Boxplot of 'diameter' in data_1 vs. 'condition_code' classes


plt.figure(figsize = (10, 6))
sns.boxplot(x = 'condition_code', y = 'diameter', data = data_1)
plt.yscale('log')
plt.show()
```



```
In [41]: # Boxplot confirms that most of the diameter values are small -- between 0 and 10 km
# Everything above 10 km is considered outliers
```


```
In [42]: # Boxplot of 'diameter' in data_1 vs. 'neo' classes
```

```
plt.figure(figsize = (10, 6))  
sns.boxplot(x = 'neo', y = 'diameter', data = data_1)  
plt.yscale('log')  
plt.show()
```



In [43]: *# Boxplot of 'diameter' in data_1 vs. 'pha' classes*

```
plt.figure(figsize = (10, 6))
sns.boxplot(x = 'pha', y = 'diameter', data = data_1)
plt.yscale('log')
plt.show()
```



In [44]: *# The distribution of values by 'neo' and 'pha' classes is similar to that by 'condition_code' classes -->*
Everything greater than 10 km is considered outlier

In [45]: *# This concludes Data Processing and EDA section*

In [46]: *# 2) Apply XGBRegressor*

In [47]: *# Separate features and target from data_1 which we will use with the xgb model*

```
X_1 = data_1.iloc[:, :-1].values # all columns, but last
y_1 = data_1.iloc[:, -1].values # last column

X_2 = data_2.values # data_2 has only features
```

In [48]: *# Split X_1 and y_1 in train/test sets*

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_1, y_1, test_size = 0.2,
                                                    random_state = 0)
```

```
In [49]: # Create XGBRegressor model

from xgboost import XGBRegressor

xgb_ini = XGBRegressor(objective = 'reg:squarederror') # denote model as 'ini'
to distinguish from optimized model later on
```

```
In [50]: # Fit & predict

xgb_ini.fit(X_train, y_train)

y_pred_1a = xgb_ini.predict(X_test) # use indexing '_1a' for comparison with l
ater predictions
```


```
In [51]: # Plot predicted diameter values in ascending order
# Log10 is used in order to display well all values

plt.figure(figsize = (12, 8))

plt.scatter(np.arange(1, len(y_test) + 1), np.sort(np.log10(y_pred_1a)), s = 50, color = 'b')

plt.xlabel('Data Point #', fontsize = 20, labelpad = 15)
plt.ylabel('Log10 of Predicted Diameter', fontsize = 20, labelpad = 15)
plt.title('XGBRegressor Model Predicted Diameter Values for Test Data', fontsi
ze = 22, c = 'b', pad = 20)
plt.tick_params(labelsize = 15)
plt.show()
```

XGBRegressor Model Predicted Diameter Values for Test Data



```
In [52]: # Main observations from plot
          # 1) small portion of predicted values are smaller than 1 km -- shown as n
          #      egative values on the plot
          # 2) largest value is approximately 500 km ( $10^{2.7}$ )
          # 3) plot shows that vast majority of predicted values are less than 10 Km
          ( $10^1$ )
```



```
In [53]: # Compare predictions, y_pred_1a, to test values, y_test, using scatterplot
# create line to represent perfect fit to y_test

y_line = np.arange(int(y_test.min()) - 10, int(y_test.max()) + 10)

# set axes limits - adjust if necessary
x_min = 0
x_max = y_test.max() + 100
d_x = 100

y_min = 0
y_max = y_test.max() + 100
d_y = 100

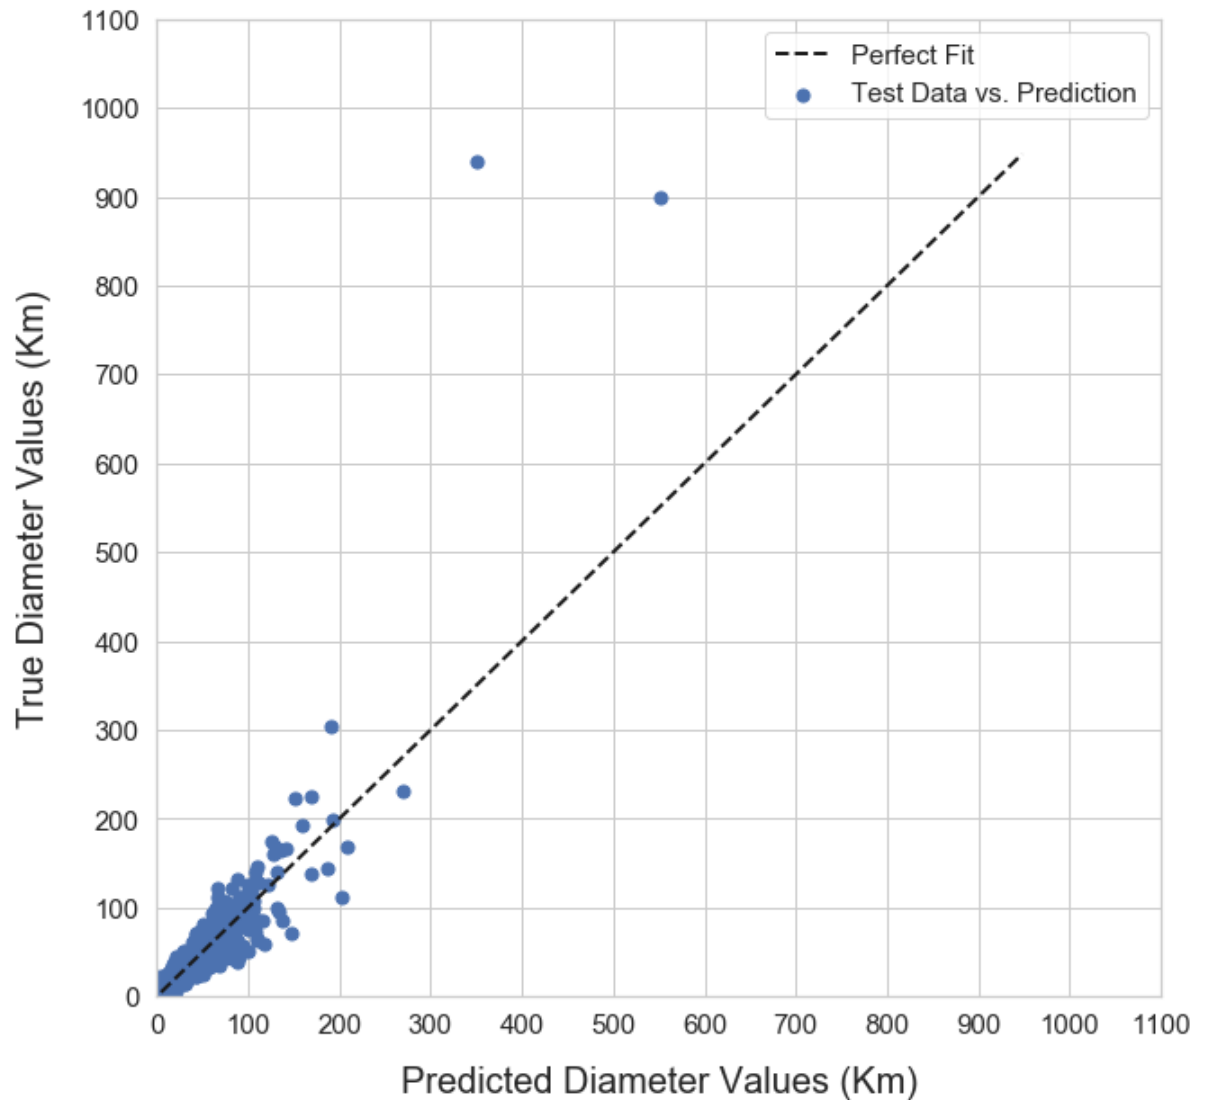
plt.figure(figsize = (10, 10))
ax = plt.axes()

ax.set_xlim(x_min, x_max)
ax.set_xticks(np.arange(x_min, x_max + d_x, d_x))

ax.set_ylim(y_min, y_max)
ax.set_yticks(np.arange(y_min, y_max + d_y, d_y))

plt.scatter(y_pred_1a, y_test, s = 50, c = 'b', label = 'Test Data vs. Prediction')
plt.plot(y_line, y_line, 'k--', lw = 2, label = 'Perfect Fit')
plt.xlabel('Predicted Diameter Values (Km)', fontsize = 20, labelpad = 15)
plt.ylabel('True Diameter Values (Km)', fontsize = 20, labelpad = 15)
plt.title('XGBRegressor Model Prediction', fontsize = 22, c = 'b', pad = 20)
plt.legend(fontsize = 15)
plt.tick_params(labelsize = 15)
plt.show()
```

XGBRegressor Model Prediction



```
In [54]: # Except for two "outliers", predictions are closely grouped around the perfect fit line
# Note: this will be discussed again later, but we would like to mention it here regarding the perceived "outliers".
# Perhaps the only limitation of XGBoost is that its predictions are capped by the data used for training
# From the scatter plot of all diameter values in the EDA section and the current plot,
# it is clear that the training data contained only points with diameters smaller than 600 km.
# That's why the predictions with the test data could not capture well the two points with diameter greater than 800 km
```

```
In [55]: # Examine model predictions in a more quantitative way --> view statistics of residuals
```

In [56]: *# Get residuals*

```
residuals_1a = y_test - y_pred_1a  
residuals_1a
```

Out[56]: array([-0.91383727, -1.68171666, -0.20006578, ..., 0.16652586,
 -0.47324387, 0.01414502])

In [57]: *# Get residuals mean and standard deviation, sigma*

```
print("Residuals_ini Mean:", round(residuals_1a.mean(),4))  
print("Residuals_ini Sigma:", round(residuals_1a.std(),4))
```

Residuals_ini Mean: 0.0187

Residuals_ini Sigma: 4.9317

In [58]: *# Mean is close to zero and sigma is small compared to the test diameter value
s, y_test -->
indicates good model accuracy*

```
In [59]: # Examine further: plot the histograms of the residuals -->
         # for better visualization plot histogram only for values within two sigma
         # s from the mean (~ 95% of all data points)

         # Set axes limits - adjust if necessary
         x_min = -10
         x_max = 10
         d_x = 2


         y_min = 0
         y_max = 4000
         d_y = 500

         plt.figure(figsize = (10, 10))
         ax = plt.axes()

         ax.set_xlim(x_min, x_max)
         ax.set_xticks(np.arange(x_min, x_max + d_x, d_x))

         ax.set_ylim(y_min, y_max)
         ax.set_yticks(np.arange(y_min, y_max + d_y, d_y))

         plt.hist(residuals_1a, bins = 2000, color = 'b')
         plt.xlabel('Residual Values', fontsize = 20, labelpad = 15)
         plt.ylabel('Count', fontsize = 20, labelpad = 15)
         plt.title('Histogram of Residuals', fontsize = 22, c = 'b', pad = 20)
         plt.tick_params(labelsize = 15)
         plt.show()
```



```
In [60]: # Visually the histogram appears close to normal distribution -->  
         # it is skewed slightly towards positive values which means that model is  
         # slightly underevaluating  
         # this can also be seen from the scatter plot of y_test vs y_pred
```

```
In [61]: # Predict the asteroid diameter values for the asteroids with unknown diameter  
         r, data_2, using model xgb_ini  
  
y_pred_2a = xgb_ini.predict(X_2)
```

```
In [62]: # Examine properties of predicted values by creating few simple plots
```

```
In [63]: # Plot predicted diameter values in ascending order


# Log10 is used in order to display well all values

plt.figure(figsize = (12, 8))

plt.scatter(np.arange(1, len(X_2) + 1), np.sort(np.log10(y_pred_2a)), s = 50, c
            = 'b')

plt.xlabel('Data Point #', fontsize = 20, labelpad = 15)
plt.ylabel('Log10 of Predicted Diameter', fontsize = 20, labelpad = 15)
plt.title('XGBRegressor Model Predicted Diameter Values for Data_2', fontsize
          = 22, c = 'b', pad = 20)
plt.tick_params(labelsize = 15)
plt.show()
```

XGBRegressor Model Predicted Diameter Values for Data_2




```
In [64]: # The range of the predicted unknown diameter values is similar to that of the
          # predicted test values
```

```
In [65]: # Plot histogram of predicted diameter values
# For better visualization, limit histogram to diameter values smaller than 20
Km

plt.figure(figsize = (10, 10))

plt.hist(y_pred_2a[y_pred_2a < 20], bins = 100, color = 'b') # limit prediction values to 20 Km
plt.xlabel('Predicted Diameter (Km)', fontsize = 20, labelpad = 15)
plt.ylabel('Histogram', fontsize = 20, labelpad = 15)
plt.title('Histogram of the Predicted (Unknown) Diameter', fontsize = 22, c = 'b', pad = 20)
plt.tick_params(labelsize = 15)
plt.show()
```

Histogram of the Predicted (Unknown) Diameter




```
In [66]: # Majority of predicted diameter values are capped by 12 km --> use this number as upper limit for histogram

plt.figure(figsize = (10, 10))

plt.hist(y_pred_2a[y_pred_2a < 12], bins = 100, color = 'b') # limit predicted values to 12 Km
plt.xlabel('Predicted Diameter (Km)', fontsize = 20, labelpad = 15)
plt.ylabel('Count', fontsize = 20, labelpad = 15)
plt.title('Histogram of the Predicted (Unknown) Diameter', fontsize = 22, c = 'b', pad = 20)
plt.tick_params(labelsize = 15)
plt.show()
```

Histogram of the Predicted (Unknown) Diameter



```
In [67]: # Predicted values for unknown diameter (data_2) have Poisson-like distribution with most of the values between 1.5 and 4 km
```



```

In [68]: # Examine if distribution is similar to distribution of known diameter values
-->
        # for adequate comparison set x-axis limit to 20 km

# set axes limits - adjust if necessary
x_min = 0
x_max = 20
d_x = 2


fig, axes = plt.subplots(1, 2, sharey = False, figsize=(16,8))

# known diameter values
axes[0].hist(y_1, bins = 2000, color = 'b')
axes[0].set_title('Histogram of Known Diameters - Data_1', fontsize = 20, c =
'b', pad = 20)
axes[0].set_xlabel('Diameter Value (km)', fontsize = 18, labelpad = 15)
axes[0].set_ylabel('Count', fontsize = 18, labelpad = 15)
axes[0].set_xlim(x_min, x_max)
axes[0].set_xticks(np.arange(x_min, x_max + d_x, d_x))
axes[0].tick_params(labelsize = 14)

# predicted unknown diameter values
axes[1].hist(y_pred_2a, bins = 1500, color = 'b')
axes[1].set_title('Histogram of Predicted Diameters - Data_2', fontsize = 20,
c = 'b', pad = 20)
axes[1].set_xlabel('Diameter Value (km)', fontsize = 18, labelpad = 15)
axes[1].set_xlim(x_min, x_max)
axes[1].set_xticks(np.arange(x_min, x_max + d_x, d_x))
axes[1].tick_params(labelsize = 14)

plt.show()

```



```
In [69]: # It appears that predicted diameter values have significantly narrower distribution encompassing smaller values

# Couple of comments regarding this observation:

# 1) We do not know how the data with known and unknown asteroid diameter were collected -->
# Thus, comparison between these two histograms may not be fully justified

# 2) However, assuming that the two sets of data are derived from the same astronomical observation and taking into account
# that the number of observation for asteroids with unknown diameter is ~ 5 times greater than that of known diameter
# one would expect that the predicted values should have similar or even wider distribution -->
# Based on item #2, the question whether an optimized model would provide different results arises
```

```
In [70]: # 3) Model optimization via hyperparameter tuning
```

```
In [71]: # For tuning use 'max_depth', 'min_child_weight', 'gamma', 'n_estimators', 'learning_rate', and 'subsample' with Randomized Search

# Note: It is important to select appropriate ranges for each of these parameters

grid_random = {'max_depth': [3, 6, 10, 20],
               'min_child_weight': np.arange(1, 10, 1),
               'gamma': np.arange(0, 10, 1),
               'n_estimators': [50, 100, 150],
               'learning_rate': [0.001, 0.01, 0.1, 0.2],
               'subsample': np.arange(0.5, 1.0, 0.1)}

from sklearn.model_selection import RandomizedSearchCV

xgb = XGBRegressor(objective = 'reg:squarederror')

xgb_random = RandomizedSearchCV(estimator = xgb, param_distributions = grid_random, n_iter = 100, cv = 5, verbose = 2,
                                random_state = 42, n_jobs = -1)
```

In [72]: *# Fit xgb_random with X_train, y_train (using same data as with xgb_ini)*

```
xgb_random.fit(X_train, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 17 tasks      | elapsed: 1.8min
[Parallel(n_jobs=-1)]: Done 138 tasks     | elapsed: 9.2min
[Parallel(n_jobs=-1)]: Done 341 tasks     | elapsed: 25.5min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 38.8min finished
```

```
Out[72]: RandomizedSearchCV(cv=5, error_score=nan,
                             estimator=XGBRegressor(base_score=0.5, booster='gbtree',
                                                      colsample_bylevel=1,
                                                      colsample_bynode=1,
                                                      colsample_bytree=1, gamma=0,
                                                      importance_type='gain',
                                                      learning_rate=0.1, max_delta_step=
0,
                                                      max_depth=3, min_child_weight=1,
                                                      missing=None, n_estimators=100,
                                                      n_jobs=1, nthread=None,
                                                      objective='reg:squarederror',
                                                      random_state=0, reg_...
                             iid='deprecated', n_iter=100, n_jobs=-1,
                             param_distributions={'gamma': array([0, 1, 2, 3, 4, 5, 6,
7, 8, 9]),
                                                  'learning_rate': [0.001, 0.01, 0.1,
0.2],
                                                  'max_depth': [3, 6, 10, 20],
                                                  'min_child_weight': array([1, 2, 3,
4, 5, 6, 7, 8, 9]),
                                                  'n_estimators': [50, 100, 150],
                                                  'subsample': array([0.5, 0.6, 0.7, 0.
8, 0.9])}),
                             pre_dispatch='2*n_jobs', random_state=42, refit=True,
                             return_train_score=False, scoring=None, verbose=2)
```

In [73]: *# Print best score and best model parameters*

```
print("Best score: %f with %s" % (xgb_random.best_score_, xgb_random.best_params_))
```

```
Best score: 0.866256 with {'subsample': 0.8999999999999999, 'n_estimators': 1
00, 'min_child_weight': 6, 'max_depth': 3, 'learning_rate': 0.1, 'gamma': 5}
```

In [74]: *# Get best_estimator_*

```
xgb_opt = xgb_random.best_estimator_
```

In [75]: *# Use xgb_opt to predict for X_test data and compare with the true values, y_test*

```
y_pred_1b = xgb_opt.predict(X_test)
```

```
In [76]: # Compare predictions, y_pred_1b, to test values, y_test

# create line to represent perfect fit to data test values, y_test

y_line = np.arange(int(y_test.min()) - 10, int(y_test.max()) + 10)

# set axes limits - adjust if necessary
x_min = 0
x_max = y_test.max() + 100
d_x = 100


y_min = 0
y_max = y_test.max() + 100
d_y = 100

plt.figure(figsize = (10, 10))
ax = plt.axes()

ax.set_xlim(x_min, x_max)
ax.set_xticks(np.arange(x_min, x_max + d_x, d_x))

ax.set_ylim(y_min, y_max)
ax.set_yticks(np.arange(y_min, y_max + d_y, d_y))

plt.scatter(y_pred_1b, y_test, s = 50, c = 'b', label = 'Test Data vs. Prediction')
plt.plot(y_line, y_line, 'k--', lw = 2, label = 'Perfect Fit')
plt.xlabel('Predicted Diameter Values (Km)', fontsize = 20, labelpad = 15)
plt.ylabel('True Diameter Values (Km)', fontsize = 20, labelpad = 15)
plt.title('XGBRegressor Model Prediction', fontsize = 22, c = 'b', pad = 20)
plt.legend(fontsize = 15)
plt.tick_params(labelsize = 15)
plt.show()
```



```
In [77]: # Predictions from Optimized model have similar behavior to those from Initial model
```

```
In [78]: # Get residuals
```

```
residuals_1b = y_test - y_pred_1b  
residuals_1b
```

```
Out[78]: array([-1.05671079, -1.6761534 , -0.22539156, ...,  0.15230133,  
               -0.62706057, -0.06124246])
```

In [79]: *# Compare residuals mean and standard deviation, sigma, from Initial and Optimized models*

```
print("Residuals_ini Mean:", round(residuals_1a.mean(),4))
print("Residuals_ini Sigma:", round(residuals_1a.std(),4))
print('\n')
print("Residuals_opt Mean:", round(residuals_1b.mean(),4))
print("Residuals_opt Sigma:", round(residuals_1b.std(),4))
```

```
Residuals_ini Mean: 0.0187
Residuals_ini Sigma: 4.9317
```

```
Residuals_opt Mean: 0.0292
Residuals_opt Sigma: 6.0209
```

In [80]: *# Although numbers are close, Initial model has slightly better residuals mean and sigma*
Thus, we should consider changing the number of hyperparameter, their ranges and the optimization method used -->
Perhaps, use Bayesian optimization instead of Randomized Search
This goes beyond the scope of this project and will be explored in subsequent project

```
In [81]: # Compare histograms of residuals from Initial and Optimized models --> for better comparison plot histograms on same graph

# Set axes limits - adjust if necessary
x_min = -6
x_max = 6
d_x = 2


y_min = 0
y_max = 4000
d_y = 500

plt.figure(figsize = (10, 10))
ax = plt.axes()

ax.set_xlim(x_min, x_max)
ax.set_xticks(np.arange(x_min, x_max + d_x, d_x))

ax.set_ylim(y_min, y_max)
ax.set_yticks(np.arange(y_min, y_max + d_y, d_y))

plt.hist(residuals_1a, bins = 2000, color = 'r', label = 'Initial Model')
plt.hist(residuals_1b, bins = 2400, color = 'g', alpha = 0.7, label = 'Optimized Model')
plt.xlabel('Residual Values', fontsize = 20, labelpad = 15)
plt.ylabel('Count', fontsize = 20, labelpad = 15)
plt.title('Histogram of Residuals', fontsize = 22, c = 'b', pad = 20)
plt.tick_params(labelsize = 15)
plt.legend(fontsize = 12)
plt.show()
```



In [82]: *# Histograms are nearly identical --> difficult to make the case for one model over the other*

In [83]: *# Use xgb_opt to predict diameters for the set with unknown diameter, X_2*
`y_pred_2b = xgb_opt.predict(X_2)`



```
In [84]: # Plot predicted diameter values in ascending order
# Log10 is used in order to display well all values

plt.figure(figsize = (12, 8))

plt.scatter(np.arange(1, len(X_2) +1), np.sort(np.log10(y_pred_2b)), s = 50, c
olor = 'b')

plt.xlabel('Data Point #', fontsize = 20, labelpad = 15)
plt.ylabel('Log10 of Predicted Diameter', fontsize = 20, labelpad = 15)
plt.title('Predicted Diameter Values for Data_2 - Optimized Model', fontsize =
22, c = 'b', pad = 20)
plt.tick_params(labelsize = 15)
plt.show()
```

Predicted Diameter Values for Data_2 - Optimized Model



```
In [85]: # Compare histograms of predicted diameters from Initial and Optimized models
-->
        # for better comparison plot histograms on same graph

# Set axes limits - adjust if necessary
x_min = 0
x_max = 8
d_x = 1

y_min = 0
y_max = 80000
d_y = 5000

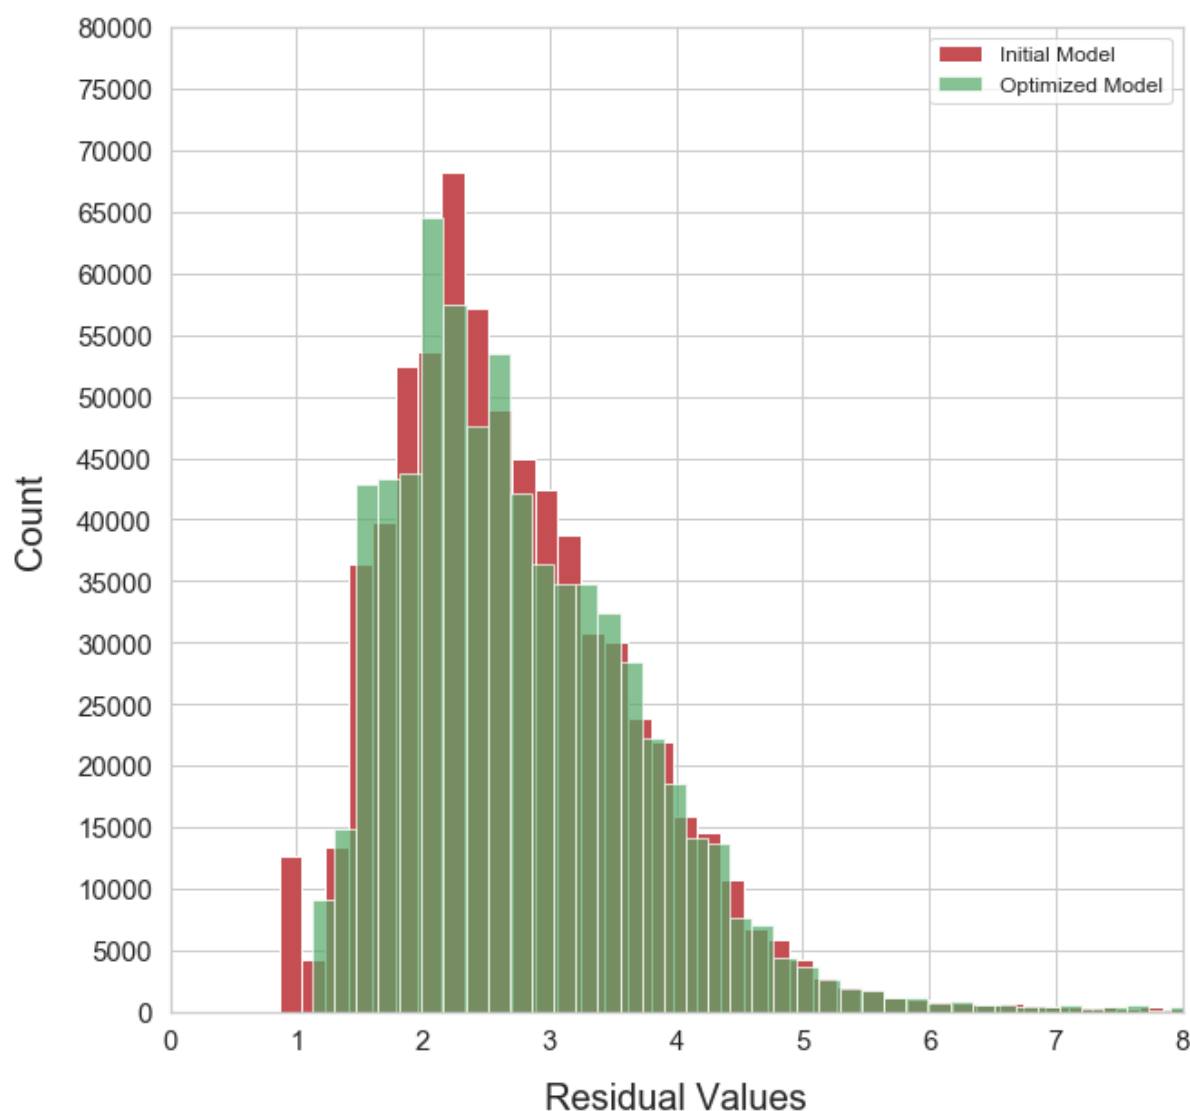
plt.figure(figsize = (10, 10))
ax = plt.axes()

ax.set_xlim(x_min, x_max)
ax.set_xticks(np.arange(x_min, x_max + d_x, d_x))

ax.set_ylim(y_min, y_max)
ax.set_yticks(np.arange(y_min, y_max + d_y, d_y))

plt.hist(y_pred_2a, bins = 3000, color = 'r', label = 'Initial Model')
plt.hist(y_pred_2b, bins = 2000, color = 'g', alpha = 0.7, label = 'Optimized
Model')
plt.xlabel('Residual Values', fontsize = 20, labelpad = 15)
plt.ylabel('Count', fontsize = 20, labelpad = 15)
plt.title('Histogram of Predicted Unknown Diameters', fontsize = 22, c = 'b',
pad = 20)
plt.tick_params(labelsize = 15)
plt.legend(fontsize = 12)
plt.show()
```

Histogram of Predicted Unknown Diameters



In [86]: *# As with the residuals, the predicted unknown diameters distributions are nearly identical for both models*
Indicates that difference in distributions between known and predicted unknown diameters is not an issue of model optimization
Because of the slightly better residual parameters we will use Initial model for final predictions

In [87]: *# 4) Final Predictions of Unknown Diameter*

In [88]: *# Train Initial model with entire X_1, y_1 dataset and use that model to predict unknown diameters for X_2*
Reason: XGBoost predictions are capped by the training data and the max values might be missing there (as in our case)
By using the entire 'data_1' dataset model should achieve the most accurate predictions possible

```
xgb_ini.fit(X_1, y_1)
y_pred_fin = xgb_ini.predict(X_2)
```


```
In [89]: # Plot final predicted diameter values in ascending order
# Log10 is used in order to display well all values

plt.figure(figsize = (12, 8))

plt.scatter(np.arange(1, len(X_2) + 1), np.sort(np.log10(y_pred_fin)), s = 50,
color = 'b')

plt.xlabel('Data Point #', fontsize = 20, labelpad = 15)
plt.ylabel('Log10 of Predicted Diameter', fontsize = 20, labelpad = 15)
plt.title('Final Predicted Diameter Values for Data_2', fontsize = 22, c = 'b',
, pad = 20)
plt.tick_params(labelsize = 15)
plt.show()
```

Final Predicted Diameter Values for Data_2



```
In [90]: # As expected, the maximum predicted value is now capped at 1000 km (10 ** 3)
```

```
In [91]: # Finally, combine the predicted diameter values with features from data_2 to
         complete the data as our final deliverable

data_2.head(10)
```

Out[91]:

	a	e	i	om	w	q	ad	per_y	data
681	2.654040	0.171983	11.505648	190.799958	104.993826	2.197591	3.110489	4.323837	40087.
698	2.610998	0.410284	15.299180	242.551766	91.399514	1.539746	3.682249	4.219081	42540.
718	2.638780	0.546301	11.564845	183.887287	156.163668	1.197212	4.080348	4.286601	39478.
729	2.243362	0.177505	4.234895	95.073806	123.549777	1.845154	2.641570	3.360139	39112.
842	2.279598	0.209766	7.997717	4.071363	316.957206	1.801415	2.757780	3.441878	37579.
961	2.908998	0.097329	2.602636	145.481660	223.473847	2.625868	3.192128	4.961619	37450.
984	2.299979	0.277462	4.056565	290.307048	59.553605	1.661822	2.938137	3.488142	35366.
1008	2.625175	0.455500	15.769676	229.461495	186.428747	1.429408	3.820942	4.253492	34990.
1010	2.391976	0.350864	5.494744	132.525452	353.279770	1.552718	3.231235	3.699504	34919.
1064	2.360276	0.297141	8.362855	330.324142	353.652287	1.658942	3.061610	3.626205	33882.

```
In [92]: # Reset data_2 indices before adding the predicted diameter values

data_2 = data_2.reset_index(drop = True)

data_2.head(10)
```

Out[92]:

	a	e	i	om	w	q	ad	per_y	data_ar
0	2.654040	0.171983	11.505648	190.799958	104.993826	2.197591	3.110489	4.323837	40087.
1	2.610998	0.410284	15.299180	242.551766	91.399514	1.539746	3.682249	4.219081	42540.
2	2.638780	0.546301	11.564845	183.887287	156.163668	1.197212	4.080348	4.286601	39478.
3	2.243362	0.177505	4.234895	95.073806	123.549777	1.845154	2.641570	3.360139	39112.
4	2.279598	0.209766	7.997717	4.071363	316.957206	1.801415	2.757780	3.441878	37579.
5	2.908998	0.097329	2.602636	145.481660	223.473847	2.625868	3.192128	4.961619	37450.
6	2.299979	0.277462	4.056565	290.307048	59.553605	1.661822	2.938137	3.488142	35366.
7	2.625175	0.455500	15.769676	229.461495	186.428747	1.429408	3.820942	4.253492	34990.
8	2.391976	0.350864	5.494744	132.525452	353.279770	1.552718	3.231235	3.699504	34919.
9	2.360276	0.297141	8.362855	330.324142	353.652287	1.658942	3.061610	3.626205	33882.

In [93]: *# Transform y_pred_1b array into series with name 'diameter'*

```
y_pred_fin = pd.Series(y_pred_fin, name = 'diameter')
y_pred_fin.head(10)
```

Out[93]:

0	12.809157
1	15.107341
2	7.265679
3	6.928519
4	7.174271
5	16.591328
6	8.169825
7	7.139175
8	8.464670
9	11.164869

Name: diameter, dtype: float32

In [94]: *# Combine features with predicted diameter values*

```
data_2 = pd.concat([data_2, y_pred_fin], axis = 1)
data_2.head(10)
```

Out[94]:

	a	e	i	om	w	q	ad	per_y	data_ar
0	2.654040	0.171983	11.505648	190.799958	104.993826	2.197591	3.110489	4.323837	40087.
1	2.610998	0.410284	15.299180	242.551766	91.399514	1.539746	3.682249	4.219081	42540.
2	2.638780	0.546301	11.564845	183.887287	156.163668	1.197212	4.080348	4.286601	39478.
3	2.243362	0.177505	4.234895	95.073806	123.549777	1.845154	2.641570	3.360139	39112.
4	2.279598	0.209766	7.997717	4.071363	316.957206	1.801415	2.757780	3.441878	37579.
5	2.908998	0.097329	2.602636	145.481660	223.473847	2.625868	3.192128	4.961619	37450.
6	2.299979	0.277462	4.056565	290.307048	59.553605	1.661822	2.938137	3.488142	35366.
7	2.625175	0.455500	15.769676	229.461495	186.428747	1.429408	3.820942	4.253492	34990.
8	2.391976	0.350864	5.494744	132.525452	353.279770	1.552718	3.231235	3.699504	34919.
9	2.360276	0.297141	8.362855	330.324142	353.652287	1.658942	3.061610	3.626205	33882.

In [95]: *# Data is complete, predicted asteroid diameter values are included, and we have accomplished project's objective*