

```
In [1]: # Using ResNet-18 for image classification
# Data:
#   # cinic10 --> https://www.kaggle.com/mengcius/cinic10
#   # a subset of cinic10 containing only four classes - airplane, automobile,
#   # ship, and truck - is used
#   # files from cinic10 'train' and 'val' folders are combined into one t
#   # o provide 18000 training images from each class
#   # test folders contain 9000 images from each class

# Note: here GPU processing is used --> running times with CPU would be greater
```

```
In [2]: # import libraries

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms, models
from torchvision.utils import make_grid
import os

import numpy as np
import pandas as pd
import seaborn as sns # for heatmaps
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style = "whitegrid", font_scale = 1.2)

# ignore non-critical warnings
import warnings
warnings.filterwarnings("ignore")

# check if GPU computing with CUDA is available and set the device accordingly
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
In [3]: # define transforms

cinic_mean = [0.4789, 0.4723, 0.4305] # from cinic10 'Readme' file
cinic_std = [0.2421, 0.2383, 0.2587] # from cinic10 'Readme' file

# we are using the same transforms for both train and test images
img_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean = cinic_mean, std = cinic_std)
])
```

```
In [4]: # prepare train and test sets, loaders

# define root directory path
root = 'datasets/cinic10_classes4'

# batch size
b_size = 16

train_data = datasets.ImageFolder(os.path.join(root, 'train'), transform = img
_transform)
test_data = datasets.ImageFolder(os.path.join(root, 'test'), transform = img_t
ransform)

torch.manual_seed(42)
train_loader = DataLoader(train_data, batch_size = b_size, shuffle = True)
test_loader = DataLoader(test_data, batch_size = b_size, shuffle = True)

class_names = train_data.classes

print(class_names)
print(f'Training images available: {len(train_data)}')
print(f'Testing images available: {len(test_data)}')

['airplane', 'automobile', 'ship', 'truck']
Training images available: 72000
Testing images available: 36000
```

```
In [5]: # display a batch of images --> grab first batch of 16 images from train data

for images, labels in train_loader:
    break

images.shape
```

```
Out[5]: torch.Size([16, 3, 32, 32])
```

```
In [6]: # 16 images in one batch, 3 color channels (R-G-B), 32 x 32 pixels
```

```

In [7]: # show images and their classes

# print labels
print('Label:', labels.numpy())
print('Class:', *np.array([class_names[i] for i in labels]))

im = make_grid(images) # the default nrow is 8

# inverse normalize the images

inv_normalize = transforms.Normalize(
    mean=[-0.4789/0.2421, -0.4723/0.2383, -0.4305/0.2587],
    std=[1/0.2421, 1/0.2383, 1/0.2587]
)
im_inv = inv_normalize(im)

# print images
plt.figure(figsize=(16,6))
plt.imshow(np.transpose(im_inv.numpy(), (1, 2, 0)));

```

Label: [1 1 1 1 1 3 2 2 1 1 1 1 0 0 3 0]

Class: automobile automobile automobile automobile automobile truck ship ship  
 automobile automobile automobile automobile airplane airplane truck airplane



```

In [8]: # everything works fine!

```

```

In [9]: # select resnet18 model from torchvision models
ResNet18model = models.resnet18()

```

```
In [10]: # display structure of resnet18  
ResNet18model
```

```

Out[10]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), b
ias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running
stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_m
ode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
    )
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
  )
)

```

```

        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=

```

```
(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
  )
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=1000, bias=True)
)
```

```
In [11]: # modify last layer of model to have 4 outputs --> we work with four classes only
torch.manual_seed(42)

ResNet18model.fc = nn.Sequential(nn.Linear(512, 4),
                                  nn.LogSoftmax(dim=1))

ResNet18model.to(device) # set device
```



```

Out[11]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), b
ias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running
stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_m
ode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn
ing_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
    )
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_run
ning_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1), bias=False)

```

```

        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=

```

```
(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Sequential(
    (0): Linear(in_features=512, out_features=4, bias=True)
    (1): LogSoftmax()
  )
)
```

In [12]: *# Last layer of model has been modified to meet our needs*

```
In [13]: # count model parameters

def count_parameters(model):
    params = [p.numel() for p in model.parameters() if p.requires_grad]
    for item in params:
        print(f'{item:>6}')
    print(f'_____ \n{sum(params):>6}')

count_parameters(ResNet18model)
```

9408  
64  
64  
36864  
64  
64  
36864  
64  
64  
36864  
64  
64  
36864  
64  
64  
73728  
128  
128  
147456  
128  
128  
8192  
128  
128  
147456  
128  
128  
147456  
128  
128  
294912  
256  
256  
589824  
256  
256  
32768  
256  
256  
589824  
256  
256  
589824  
256  
256  
1179648  
512  
512  
2359296  
512  
512  
131072  
512  
512  
2359296  
512  
512

```
2359296
  512
  512
 2048
   4
-----
11178564
```

```
In [14]: # approximately 11 million parameters which is not so large of a number
```

```
In [15]: # define loss function and optimizer

# define loss function and set device
criterion = nn.CrossEntropyLoss().to(device)

# optimizer needs to be defined after device for model is selected --> model parameters with different devices are different!
optimizer = torch.optim.Adam(ResNet18model.parameters(), lr = 0.001)
```

```

In [16]: # start training

import time
start_time = time.time()

epochs = 40

# Limit number of train and test images --> optional - saves time to run with
# small numbers first to see if model works
max_trn_batch = 1000
max_tst_batch = 500

# instantiate trackers for model performance
train_losses = []
test_losses = []
train_correct = []
test_correct = []

for i in range(epochs):
    trn_corr = 0
    tst_corr = 0

    # run the training batches
    for b, (X_train, y_train) in enumerate(train_loader):

        # Limit the number of batches
        if b == max_trn_batch:
            break
        b+=1

        # apply the model
        y_pred = ResNet18model(X_train.to(device))
        loss = criterion(y_pred, y_train.to(device))

        # tally the number of correct predictions
        predicted = torch.max(y_pred.data, 1)[1]
        batch_corr = (predicted == y_train.to(device)).sum()
        trn_corr += batch_corr

        # update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # print interim results
        if b%max_trn_batch == 0:
            print(f'Epoch: {i+1:3}  Batch: {b:4}  Loss: {loss.item():10.4f}
Accuracy: {trn_corr.item()*100/(b_size*b):7.2f}% \
Time Elapsed Since Start: {(time.time() - start_time)/60:7.2f} min')

    train_losses.append(loss)
    train_correct.append(trn_corr)

    # run the test batches
    with torch.no_grad():
        for b, (X_test, y_test) in enumerate(test_loader):

```

```
# limit the number of batches
if b == max_tst_batch:
    break

# apply the model
y_val = ResNet18model(X_test.to(device))

# tally the number of correct predictions
predicted = torch.max(y_val.data, 1)[1]
tst_corr += (predicted == y_test.to(device)).sum()

loss = criterion(y_val, y_test.to(device))
test_losses.append(loss)
test_correct.append(tst_corr)
```



Epoch: 1	Batch: 1000	Loss: 0.9393	Accuracy: 49.15%	Time Elapse
d Since Start:	0.73 min			
Epoch: 2	Batch: 1000	Loss: 1.0233	Accuracy: 56.04%	Time Elapse
d Since Start:	1.53 min			
Epoch: 3	Batch: 1000	Loss: 0.9782	Accuracy: 61.43%	Time Elapse
d Since Start:	2.33 min			
Epoch: 4	Batch: 1000	Loss: 1.3356	Accuracy: 63.71%	Time Elapse
d Since Start:	3.14 min			
Epoch: 5	Batch: 1000	Loss: 0.6735	Accuracy: 66.08%	Time Elapse
d Since Start:	3.94 min			
Epoch: 6	Batch: 1000	Loss: 0.9234	Accuracy: 67.83%	Time Elapse
d Since Start:	4.82 min			
Epoch: 7	Batch: 1000	Loss: 0.6212	Accuracy: 70.58%	Time Elapse
d Since Start:	5.69 min			
Epoch: 8	Batch: 1000	Loss: 0.4811	Accuracy: 71.34%	Time Elapse
d Since Start:	6.54 min			
Epoch: 9	Batch: 1000	Loss: 0.8843	Accuracy: 72.24%	Time Elapse
d Since Start:	7.38 min			
Epoch: 10	Batch: 1000	Loss: 0.3796	Accuracy: 74.07%	Time Elapse
d Since Start:	8.21 min			
Epoch: 11	Batch: 1000	Loss: 1.0706	Accuracy: 74.71%	Time Elapse
d Since Start:	9.04 min			
Epoch: 12	Batch: 1000	Loss: 0.6988	Accuracy: 75.54%	Time Elapse
d Since Start:	9.86 min			
Epoch: 13	Batch: 1000	Loss: 0.6407	Accuracy: 76.58%	Time Elapse
d Since Start:	10.69 min			
Epoch: 14	Batch: 1000	Loss: 0.4520	Accuracy: 77.59%	Time Elapse
d Since Start:	11.52 min			
Epoch: 15	Batch: 1000	Loss: 0.5094	Accuracy: 77.51%	Time Elapse
d Since Start:	12.35 min			
Epoch: 16	Batch: 1000	Loss: 0.5058	Accuracy: 78.47%	Time Elapse
d Since Start:	13.17 min			
Epoch: 17	Batch: 1000	Loss: 0.9032	Accuracy: 78.99%	Time Elapse
d Since Start:	13.98 min			
Epoch: 18	Batch: 1000	Loss: 1.0348	Accuracy: 80.16%	Time Elapse
d Since Start:	14.79 min			
Epoch: 19	Batch: 1000	Loss: 0.6745	Accuracy: 79.69%	Time Elapse
d Since Start:	15.60 min			
Epoch: 20	Batch: 1000	Loss: 0.6603	Accuracy: 81.23%	Time Elapse
d Since Start:	16.42 min			
Epoch: 21	Batch: 1000	Loss: 0.2936	Accuracy: 82.04%	Time Elapse
d Since Start:	17.23 min			
Epoch: 22	Batch: 1000	Loss: 0.5107	Accuracy: 81.99%	Time Elapse
d Since Start:	18.04 min			
Epoch: 23	Batch: 1000	Loss: 0.9354	Accuracy: 83.46%	Time Elapse
d Since Start:	18.85 min			
Epoch: 24	Batch: 1000	Loss: 0.3793	Accuracy: 83.41%	Time Elapse
d Since Start:	19.65 min			
Epoch: 25	Batch: 1000	Loss: 0.4244	Accuracy: 84.31%	Time Elapse
d Since Start:	20.46 min			
Epoch: 26	Batch: 1000	Loss: 0.2853	Accuracy: 84.51%	Time Elapse
d Since Start:	21.30 min			
Epoch: 27	Batch: 1000	Loss: 0.9477	Accuracy: 85.10%	Time Elapse
d Since Start:	22.12 min			
Epoch: 28	Batch: 1000	Loss: 0.1592	Accuracy: 85.32%	Time Elapse
d Since Start:	22.92 min			
Epoch: 29	Batch: 1000	Loss: 0.2434	Accuracy: 86.53%	Time Elapse

```

d Since Start: 23.73 min
Epoch: 30 Batch: 1000 Loss: 0.2986 Accuracy: 86.85% Time Elapse
d Since Start: 24.53 min
Epoch: 31 Batch: 1000 Loss: 0.7745 Accuracy: 87.21% Time Elapse
d Since Start: 25.34 min
Epoch: 32 Batch: 1000 Loss: 0.3200 Accuracy: 87.58% Time Elapse
d Since Start: 26.14 min
Epoch: 33 Batch: 1000 Loss: 0.6090 Accuracy: 87.52% Time Elapse
d Since Start: 26.95 min
Epoch: 34 Batch: 1000 Loss: 0.1243 Accuracy: 88.46% Time Elapse
d Since Start: 27.75 min
Epoch: 35 Batch: 1000 Loss: 0.2393 Accuracy: 88.78% Time Elapse
d Since Start: 28.57 min
Epoch: 36 Batch: 1000 Loss: 0.3746 Accuracy: 89.17% Time Elapse
d Since Start: 29.38 min
Epoch: 37 Batch: 1000 Loss: 0.1953 Accuracy: 89.42% Time Elapse
d Since Start: 30.20 min
Epoch: 38 Batch: 1000 Loss: 0.0828 Accuracy: 90.14% Time Elapse
d Since Start: 31.01 min
Epoch: 39 Batch: 1000 Loss: 0.2272 Accuracy: 90.23% Time Elapse
d Since Start: 31.82 min
Epoch: 40 Batch: 1000 Loss: 0.2472 Accuracy: 90.61% Time Elapse
d Since Start: 32.62 min

```

In [17]: *# save the trained model*


```
torch.save(ResNet18model.state_dict(), 'resnet18cinic10cl4.pt')
```

In [18]: *# evaluate model performance*

```


# plot losses
plt.figure(figsize = (10,5))
plt.plot(train_losses, label='training loss')
plt.plot(test_losses, label='validation loss')
plt.title('Loss at the end of each epoch')
plt.legend();

```



In [19]: *# train and test losses do not converge*

In [20]: *# plot accuracy*  
`plt.figure(figsize = (10,5))`  
`plt.plot([t*100/(max_trn_batch*b_size) for t in train_correct], label='training accuracy')`  
`plt.plot([t*100/(max_tst_batch*b_size) for t in test_correct], label='validation accuracy')`  
`plt.title('Accuracy at the end of each epoch')`  
`plt.legend();`



In [21]: *# accuracy diverges as well*  
*# model achieves 90% training accuracy and ~ 77% testing (validation) accuracy*  
*# after ~ 30 epochs the validation accuracy reaches a plateau, while the training accuracy continues to increase*  
*# if the number of epochs is increased, model could reach even higher training accuracy by "learning" the training set*  
*# this, however, will not lead to significant increase in validation accuracy which is the true model accuracy*

In [22]: *# print final test accuracy*  
`print(f'Test accuracy: {test_correct[-1].item()*100/(max_tst_batch*b_size):.2f}%)` *# take the value after the last epoch*

Test accuracy: 77.00%

In [23]: *# test accuracy is 77%*

In [24]: *# compare predictions against ground truth*  
*# for better visualization we will use heatmap*

```
In [25]: # use all test images
test_load_all = DataLoader(test_data, batch_size = 36000, shuffle = False)
print(f'Testing images available: {len(test_data)}')

Testing images available: 36000
```

```

In [26]: # make predictions with all test images and show confusion matrix results as a
          heat map

          # import confusion matrix
          from sklearn.metrics import confusion_matrix

          with torch.no_grad():
              correct = 0
              for X_test, y_test in test_loader:
                  y_val = ResNet18model(X_test.to(device))
                  predicted = torch.max(y_val,1)[1]
                  correct += (predicted == y_test.to(device)).sum()

          # convert results to CPU tensors to be able to use them as numpy arrays!
          device = torch.device("cpu")
          y_test = y_test.to(device)
          predicted = predicted.to(device)

          # create heat map from confusion matrix and plot
          arr = confusion_matrix(y_test.view(-1), predicted.view(-1))
          df_cm = pd.DataFrame(arr, class_names, class_names)
          plt.figure(figsize = (10,8))
          sns.heatmap(df_cm, annot = True, fmt = "d", cmap = 'BuGn')
          plt.xlabel("Prediction")
          plt.ylabel("Label (ground truth)")
          plt.show()

```



```
In [27]: # Confusion matrix shows:
          # class 'airplane' has highest prediction accuracy of 85.4% (7690/9000)
          # most often mistaken for 'ship'
          # class 'ship' is close second with prediction accuracy of 84.7% (7622/9000)
          # most often mistaken for 'airplane'
          # class 'automobile' is third in prediction accuracy with 74.3% (6684/9000)
          # most often mistaken for 'truck'
          # class 'truck' has lowest prediction accuracy of 66.4% (5980/9000)
          # most often mistaken for 'automobile'
          # the corresponding classes causing highest confusion are natural to expect
          # due to the similarities between these classes

          # In conclusion:
          # overall model accuracy is 77.7% (27976/36000) which is well above random
          # guess (25%)
          # given the similarity between the objects selected here and the small image
          # resolution we consider this to be good accuracy
```